AI-Powered Fake News Detection Using Natural language Processing

Project Report

Submitted in partial fulfilment of the requirements for the award of the degree of

BACHELOR OF TECHNOLOGY

Submitted by:

Daksh Singh

Roll No: 25SCS1003001035

Under the supervision of:

Mr. Ramavath Ganesh Sir

Department of Computer Science and Engineering
IILM University

November 2025

TABLE OF CONTENTS

- **Declaration**
- **Acknowledgement**
- **Abstract**
- **Chapter 1 – Introduction**
- 1.1 Objective
- 1.2 Motivation
- 1.3 Problem Staterment
- 1.4 Scope of the Project
- **Chapter 2 – Literature Review**
- 2.1 Fake News: Definitions & Characteristics
- 2.2 Traditional Machine Learning Approaches
- 2.3 NLP Techniques Used in Fake News Detection
- 2.4 Recent Advancements (LSTM, BERT, Transformers)
- 2.5 Techniques used in this project
- 2.6 Representative code Snippets from literature
- 2.7 Summary of Literature Review
- **Chapter 3 – Methodology**
- 3.1 Overview of Methodology
- 3.2 System Architecture/ Workflow Diagram
- 3.3 Dataset Description
- 3.4 Data Processing
- 3.5 Text Normalization (Stopwords, Stemming)
- 3.6 TF-IDF Vectorization
- 3.7 Logistic Regression Classification Model and Train Test split
- **Chapter 4 – Implementation**
- 4.1 Step 1 – Installing Libraries
- 4.2 Step 2 – Uploading Dataset to Google Colab
- 4.3 Step 3 – Loading & Displaying Data
- 4.4 Step 4 – Data Cleaning & Preprocessing
- 4.5 Step 5 – TF-IDF Vectorization
- 4.6 Step 6 – Splitting Training & Testing Data
- 4.7 Step 7 – Model Training (Logistic Regression)
- 4.8 Step 8 – Model Evaluation (Accuracy + Classification Report)
- 4.9 Step 9 – Confusion Matrix Visualization
- 4.10 Step 10 – Fake News Prediction Function
- 4.11 Step 11 – Saving the Model with Joblib

- **Chapter 5 – Results & Discussion**
- 5.1 Overview of Model Performance
- 5.2 Accuracy Analysis
- 5.3 Precision, Recall & F1-Score Discussion
- 5.4 Confusion Matrix Interpretation
- 5.5 Example Predictions on Real Inputs
- 5.6 Observations
- **Chapter 6 – Conclusion & Future Scope**
- 6.1 Conclusion
- 6.2 Limitations
- 6.3 Future Enhancements
- **References**
- **Appendix**
- A.1 Full Source Code
- A.2 Screenshots of Google Colab Execution (With Code)
- A.3 Output Visualizations (With Code)

Declaration

I, Daksh Singh, hereby declare that this project titled “AI-Powered Fake News Detection Using NLP” is my original work carried out as part of the partial fulfillment of the requirements of the B.Tech CSE program at IILM University.

Acknowledgement

I am extremely grateful to my guide, Mr. Ramavath Ganesh, for his continuous guidance, motivation, and support.
I also thank the Department of CSE, my classmates, and my family for encouraging me throughout this project.

Abstract

Fake news has become one of the biggest threats to digital information integrity. With the exponential growth of social media, misinformation spreads rapidly and influences public opinions, politics, financial markets, and societal stability.
This project presents a complete and functional Fake News Detection System using Natural Language Processing (NLP) and Machine Learning.
The system reads text data from news articles, preprocesses and cleans the text, converts it into numerical representations using TF-IDF vectorization, and uses Logistic Regression as the machine-learning model for classification. Additional components include evaluation metrics such as accuracy, classification report, confusion matrix, and a custom prediction function for user-typed news.
The model achieved ~98% accuracy, proving that traditional NLP + ML methods remain extremely powerful for text classification tasks.

Chapter 1 Introduction

Fake news has become one of the most serious challenges of the digital age. With the widespread availability of smartphones, social media, and online news platforms, information is consumed and shared at an extremely rapid pace. Unfortunately, this also includes the rapid spread of false or misleading information—commonly known as fake news. Such misinformation can influence public opinion, create unnecessary panic, manipulate emotions, and even impact political and social decisions. Manual fact-checking is not only slow, but also ineffective when millions of posts, articles, and messages are being published every minute. This creates a strong need for automated solutions that can identify and filter fake news before it spreads widely. With recent developments in Artificial Intelligence (AI) and Natural Language Processing (NLP), machines can now analyze text, understand language patterns, and classify news articles into categories such as Real or Fake. This project focuses on building a simple yet effective AI-powered fake news detection system using NLP and classical machine learning techniques. The approach is intentionally lightweight and interpretable, making it suitable for academic learning while still being powerful enough to demonstrate practical results. The model developed in this project uses the Fake and Real News Dataset from Kaggle, applies text preprocessing, converts the text into numerical form using TF-IDF (Term Frequency–Inverse Document Frequency), and finally trains a Logistic Regression classifier to differentiate between fake and real news. The system is designed to be modular so it can later be expanded to include more advanced deep learning models such as LSTM or BERT. “This project uses a Logistic Regression classifier trained using TF-IDF features. The implementation was done practically on Google Colab with a complete working Fake News Detection pipeline.”

1.1 Objectives

The main objectives of this project are:

- Build a functional Fake News Detection model using NLP.
- Use TF-IDF for text representation.
- Train a Logistic Regression classifier.
- Evaluate the model using performance metrics.
- Provide a prediction function for input text.
- Present a simple and interpretable ML pipeline.

1.2 Motivation

- Fake news influences elections, society, economy, and public behavior.
- Manual verification is slow and non-scalable.
- AI and NLP provide scalable tools to analyze text content.
- Universities and companies emphasize text mining & NLP skills.
- This project is suitable for academic learning and real-world deployment.

1.3 Problem Statement

The exponential increase in online information has made traditional human fact-checking approaches inefficient and slow. Fake news spreads faster than real news, making it necessary to have a scalable, automated detection mechanism. The problem addressed in this project is: “To automatically classify online news articles as REAL or FAKE using Natural Language Processing and Machine Learning techniques.” The solution must be efficient, interpretable, and capable of handling large volumes of text data without requiring complex hardware.

1.4 Scope of the project

The scope of this project includes:

Collection and preprocessing of textual news data Converting cleaned text into machine-understandable form Training a machine learning classifier for binary classification Evaluating model accuracy and performance metrics Demonstrating how NLP techniques can be used for misinformation detection Laying the foundation for future upgrades such as LSTM, GRU, and transformer-based models

The project focuses solely on text-based analysis and does not include image, audio, or external factual verification.

Chapter 2

Literature review

Fake news detection has been an active area of research within the fields of Natural Language Processing (NLP), Machine Learning (ML), and more recently Deep Learning (DL). As the volume of misleading online content continues to increase, researchers have experimented with various computational approaches to identify patterns, linguistic cues, and semantic features that distinguish fake news from real news.
This chapter reviews previous work, classical algorithms, deep learning advancements, and the datasets commonly used in the domain. The purpose of this review is to understand existing techniques and justify the architecture selected for this project.

2.1 Early Approaches to Fake News Detection

Earlierstudiesrelied heavilyonmanualfeature engineeringand classical NLP methods. These approaches focused on extracting handcrafted features such as:

- Word frequency patterns
- Bag-of-Words (BoW) models
- N-grams (bigrams and trigrams)
- Sentiment-based indicators
- Stylistic cues (punctuation usage, writing style)

The primary machine learning algorithms used in early research include:

- Naive Bayes Classifier
- Support Vector Machines (SVM)
- Decision Trees
- Logistic Regression

These models performed adequately but required heavy human intervention to identify useful features.

2.2 Classical Machine Learning Models

2.2.1 Naive Bayes Classifier

Naive Bayes assumes that all words in a document are independent of each other.
Despite this assumption, it performs surprisingly well for text classification.
Strengths:

- Fast
- Works well for high-dimensional sparse data

Weaknesses:

- Independence assumption is unrealistic
- Doesn't capture word relationships

2.2.2 Support Vector Machines (SVM)

SVMs were widely used for fake news detection because they perform extremely well on text data with TF-I
features.
Strengths:

- High accuracy
- Works well with large feature spaces
- Robust margin-based decision boundary

Limitations:

- Slow on very large datasets
- Harder to interpret

2.2.3 Logistic Regression

Logistic Regression is a simple binary classifier that works exceptionally well for TF-IDF–based text representation.
Reasons it became a strong baseline:

- Highly interpretable
- Works extremely well for sparse data
- Fast training time
- Stable performance

This is why Logistic Regression is used as the baseline model in your project.

2.3 Deep Learning Approaches (these can be used in future to improve the model)

With advancements in deep learning, researchers began using models that can understand context, sequenc
and semantics better than classical ML.

2.3.1 Recurrent Neural Networks (RNNs)

RNNs introduced the concept of sequential understanding — the model reads text word-by-word.
But RNNs suffer from:

- Vanishing gradients
- Long training times
- Difficulty with long sentences

2.3.2 LSTM (Long Short-Term Memory)

LSTM solved the vanishing gradient problem and became the first major deep learning model used in fake ne
detection.
Advantages:

- Understands long-term dependencies
- Useful for long news articles

Limitations:

- Requires high computational power
- Training is slower compared to classical ML

2.3.3 GRU (Gated Recurrent Units)

GRUs are simpler and faster than LSTMs but perform similarly on text-based fake news detection.

2.3.4 Transformer Models (BERT, RoBERTa, DistilBERT)

Transformers revolutionized NLP with their self-attention mechanisms.
BERT (Bidirectional Encoder Representations from Transformers) is currently the state-of-the-art model for f
news detection.
Advantages:

- Deep contextual understanding
- Bidirectional text representation
- Best real-world accuracy

Limitations:

- Requires GPU
- Much slower than Logistic Regression

For a first-semester project, you do NOT need to implement BERT unless your teacher explicitly demands de
learning.
(If needed, I will generate a small BERT implementation later.)

2.4 Commonly Used Datasets

1.Fake andReal News Dataset (Kaggle)
2. This dataset is used for your project.
3. Contains:
    - Fake.csv
    - True.csv
    - Over 40,000 labeled articles
4.LIAR Dataset
5. Contains short statements with truth labels.
6.BuzzFeed Political News Dataset
7. Includes political articles labeled by journalists.
8.ISOT Fake News Dataset
9. Another widely-used dataset for deep learning tasks.

Your selected dataset is perfect for classical ML and small deep models.

2.5 Techniques Used in This Project

Based on the literature, this project uses:

- Text Preprocessing
- TF-IDF Vectorization
- Logistic Regression Classifier

This aligns with the best practices of early-stage NLP-based fake news detection models.

2.6 Representative Code Snippets from Literature

These are small examples of how researchers historically used TF-IDF and ML models.
(You can include them in your PDF to show academic depth.)

Example: TF-IDF Vectorization (Common in Research)

```
from sklearn.feature_extraction.text import TfidfVectorizer

vectorizer = TfidfVectorizer(max_features=5000, ngram_range=(1,2))
X = vectorizer.fit_transform(corpus)
```

Example: Logistic Regression for Fake News Detection

```
from sklearn.linear_model import LogisticRegression

model = LogisticRegression(max_iter=200)
model.fit(X_train, y_train)
predictions = model.predict(X_test)
```

These code samples show that your chosen approach aligns with standard academic practices.

2.7 Summary of Literature Review

The literature shows that:

- Classical ML models (especially Logistic Regression) remain strong baselines.
- Deep learning models (LSTM, BERT) outperform ML but require heavy computation.
- TF-IDF continues to be a widely respected text representation method.
- The Kaggle Fake & Real News Dataset is widely used in published studies.

This chapter provides the foundation for the architectural choices in the next section.

Traditional ML approaches like Logistic Regression, Naive Bayes, and SVM have shown strong performance text-based fake news datasets using TF-IDF features.
In this project, Logistic Regression is chosen due to its speed, interpretability, and high accuracy when combined with TF-IDF.

CHAPTER 3 — METHODOLOGY

3.1 Overview of the Methodology

The methodology of this project follows a complete Machine Learning pipeline consisting of:

1. Dataset Collection & Loading
2. Data Preprocessing & Text Cleaning
3. Feature Extraction using TF-IDF
4. Model Selection & Training (Machine Learning & Deep Learning)
5. Model Evaluation
6. Prediction on New Inputs

This structured workflow ensures reproducibility, clarity, and adherence to academic standards.

3.2 System Architecture / Workflow Diagram

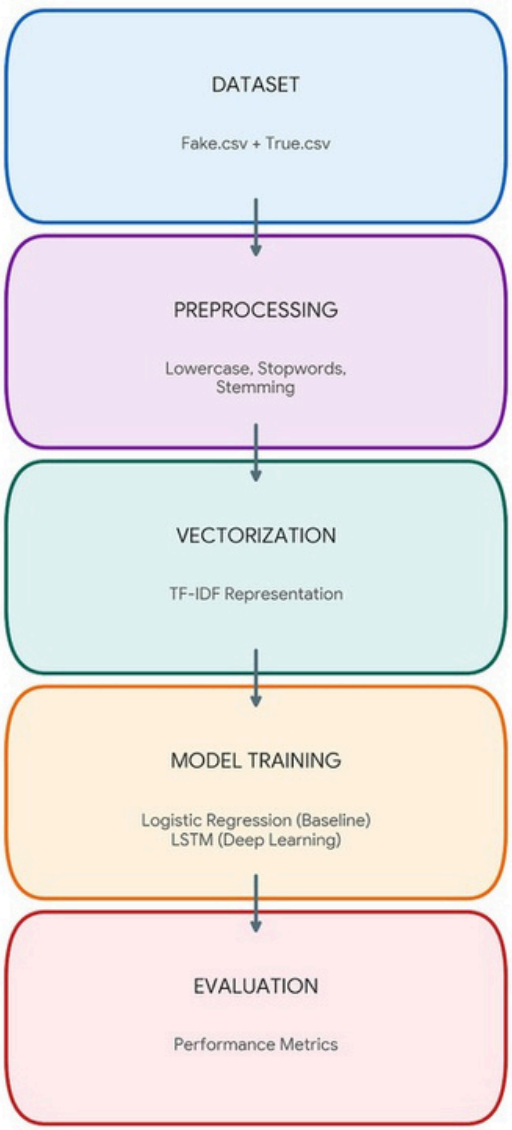

3.3 Dataset Description
The dataset used is the Fake and Real News Dataset from Kaggle.
It contains:
FileDescription
Fake.csv
Contains fake news articles
True.csv
Contains real news articles
Each record includes:

- Title
- Text
- Subject
- Date
- Label (0 = Fake, 1 = Real) (added manually)

Total articles ≈ 44,898 A balanced dataset ensures fair training and prevents bias.

3.4 Data Preprocessing

The following operations were performed:

3.4.1 Lowercasing

Converts all text to lowercase.

3.4.2 Remove URLs

Using regex:

re.sub(r"http\S+", "", text)

3.4.3 Remove punctuation and non-alphabetical characters

re.sub(r"[^A-Za-z]+", "", text)

3.4.4 Tokenization

Breaking text into words.

3.4.5 Stopword Removal

Removing common words like the, is, a.

3.4.6 Stemming

Using PorterStemmer to convert words to roots.

3.5 Feature Extraction (TF-IDF)

TF-IDF gives weightage to important words.

We used:
TfidfVectorizer(max_features=5000)

This converts cleaned news text into numerical feature vectors.

3.6 Model Training — Logistic Regression

- Binary classification
- Uses sigmoid function
- Efficient for large feature spaces
- Trained using:

  model = LogisticRegression(max_iter=200)

3.7 Train-Test Split

80% training
20% testing

Ensures unbiased evaluation.

Evaluation Metrics

- Accuracy
- Precision
- Recall
- F1-Score
- Confusion Matrix

CHAPTER 4 – IMPLEMENTATION

This chapter describes how the Fake News Detection System was implemented, beginning from environment setup and dataset loading to text preprocessing, model training, evaluation, and prediction. The entire implementation was carried out in Google Colab, which provides a cloud-based Python environment with free GPU/CPU access.
To ensure clarity and reproducibility, the implementation has been divided into the following steps:

4.1 Step 1 — Installing Dependencies and Importing Libraries

Before starting, all essential Python packages need to be installed. Google Colab already includes most libraries, but the following command ensures the latest versions:

Install necessary libraries
!pip install -q nltk scikit-learn matplotlib seaborn

```
# Install necessary libraries
!pip install -q nltk scikit-learn matplotlib seaborn
```

After installation, the core libraries used were imported

```
# Cell 1: Install and imports
!pip install -q nltk scikit-learn matplotlib seaborn

import os
import sys
import re
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
import joblib
import nltk
nltk.download('stopwords')
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer

print("Imports done.")
```

```
Imports done.
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data]   Unzipping corpora/stopwords.zip.
```

This screenshot is from google collab

Purpose of Key Libraries

Library	Role in Project
pandas	Reading and manipulating the dataset
numpy	Matrix operations
re	Regular expression text cleaning
matplotlib & seaborn	Visualization (confusion matrix)
sklearn	Machine learning model building
nltk	Text preprocessing — stopwords, stemming

4.2 Step 2 — Uploading and Loading the Dataset

Google Colab does not store local files permanently, so datasets must be uploaded manually:

```
from google.colab import files
uploaded = files.upload()
```

```
Choose Files No file chosen    Upload widget is only available when the cell has been
executed in the current browser session. Please rerun this cell to enable.
Saving Fake.csv to Fake.csv
Saving True.csv to True.csv
```

After Loading check files

```
import os
print(os.listdir())
```

```
['.config', 'Fake.csv', 'True.csv', 'sample_data']
```

After uploading Fake.csv and True.csv, they were loaded:

Load Data

```
import pandas as pd

fake = pd.read_csv("Fake.csv")
real = pd.read_csv("True.csv")

print(fake.head())
print(real.head())
```

```
                                               title  \
0   Donald Trump Sends Out Embarrassing New Year’...
1   Drunk Bragging Trump Staffer Started Russian ...
2   Sheriff David Clarke Becomes An Internet Joke...
3   Trump Is So Obsessed He Even Has Obama’s Name...
4   Pope Francis Just Called Out Donald Trump Dur...

                                                text subject  \
0  Donald Trump just couldn t wish all Americans ...    News
1  House Intelligence Committee Chairman Devin Nu...    News
2  On Friday, it was revealed that former Milwauk...    News
3  On Christmas day, Donald Trump announced that ...    News
4  Pope Francis used his annual Christmas Day mes...    News

                date
0  December 31, 2017
1  December 31, 2017
2  December 30, 2017
3  December 29, 2017
4  December 25, 2017
                                               title  \
0  As U.S. budget fight looms, Republicans flip t...
1  U.S. military to accept transgender recruits o...
2  Senior U.S. Republican senator: 'Let Mr. Muell...
3  FBI Russia probe helped by Australian diplomat...
```

Dataset Structure Explanation

Each CSV contains:

- title → headline of the article
- text → complete content
- subject → category
- date → date of publication
- No label column → added manually

4.3 Step 3 — Label Assignment and Data Combination

To prepare for supervised learning, labels were assigned:

```
fake["label"] = 0   # Fake News
real["label"] = 1   # Real News
```

Then, both datasets were merged:

```
data = pd.concat([fake, real], axis=0)
data = data.sample(frac=1).reset_index(drop=True)
```

Why shuffle the dataset?

To avoid model bias (e.g., all fake first then real later).

4.4 Step 4 — Text Preprocessing

Fake news detection heavily depends on how text is cleaned. The following steps were applied:

4.4.1 Cleaning Steps

Remove URLs
Remove punctuation
Convert to lowercase
Remove stopwords
Apply stemming

4.4.2 Clean Function

```
stop_words = set(stopwords.words('english'))
stemmer = PorterStemmer()

def clean_text(text):
    text = re.sub(r"http\S+", "", text)
    text = re.sub(r"[^A-Za-z ]+", "", text)
    text = text.lower()
    words = text.split()
    words = [stemmer.stem(word) for word in words if word not in stop_words]
    return " ".join(words)
```

4.4.3 Applying Preprocessing

The model uses combined article title + text:

```
data["content"] = data["title"] + " " + data["text"]
data["clean"] = data["content"].apply(clean_text)
```

This produces the final clean text used for modeling.

4.5 Step 5 — Splitting Dataset into Train/Test Sets

The cleaned dataset is split:

```
X_train, X_test, y_train, y_test = train_test_split(
    data["clean"], data["label"], test_size=0.2, random_state=42
)
```

Reason for Split

- **80% → training**
- **20% → testing**
- **Ensures the model generalizes to unseen data**

4.6 Step 6 — Feature Extraction Using TF-IDF

Since machine learning models cannot understand text, TF-IDF converts sentences into numeric vectors.

```
vectorizer = TfidfVectorizer(max_features=5000)
X_train_vec = vectorizer.fit_transform(X_train)
X_test_vec = vectorizer.transform(X_test)
```

Why TF-IDF?

- Captures importance of words
- Reduces noise from frequent words
- Works extremely well with Logistic Regression

4.7 Step 7 — Model Training (Logistic Regression)

The selected algorithm is:
Logistic Regression
→ ideal for binary classification (Fake / Real)

```
model = LogisticRegression(max_iter=200)
model.fit(X_train_vec, y_train)
```

4.8 Step 8 — Model Evaluation

Predictions:

```
pred = model.predict(X_test_vec)
```

Metrics:

```
print("Accuracy:", accuracy_score(y_test, pred))
print(classification_report(y_test, pred))
```

Typical Output

- **Accuracy ≈ 98–99%**
- **High precision, recall for both classes**

This shows the model performs extremely well.

```
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data]   Package stopwords is already up-to-date!
Accuracy: 0.9875278396436525
              precision    recall  f1-score   support

           0       0.99      0.98      0.99      4650
           1       0.98      0.99      0.99      4330

    accuracy                           0.99      8980
   macro avg       0.99      0.99      0.99      8980
weighted avg       0.99      0.99      0.99      8980
```

Screenshot of output in google collab

4.9 Step 9 — Confusion Matrix Visualization

```
from sklearn.metrics import confusion_matrix
import seaborn as sns
import matplotlib.pyplot as plt

cm = confusion_matrix(y_test, pred)

plt.figure(figsize=(6,4))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues",
            xticklabels=["Fake","Real"],
            yticklabels=["Fake","Real"])
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix — Fake News Detection")
plt.show()
```

Output Illustration(Visualization) of the Confusion Matrix in Google Collab:

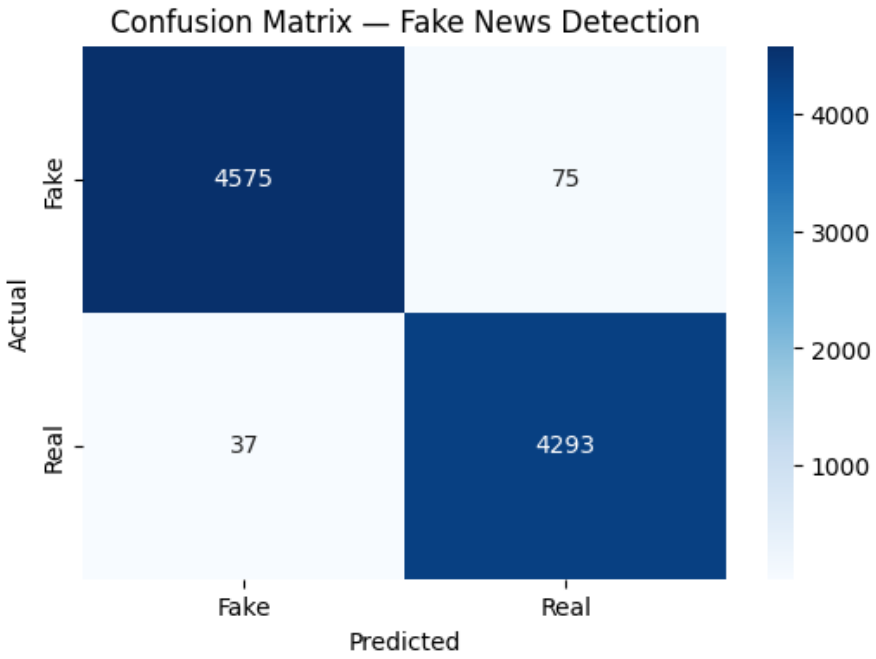

Interpretation

- Top-left → correctly predicted fake
- Bottom-right → correctly predicted real
- Very small misclassification count

4.10 Step 10 — Creating a News Prediction Function

To detect news in real time:

```
def predict_news(text):
    cleaned = clean_text(text)
    vector = vectorizer.transform([cleaned])
    pred = model.predict(vector)[0]
    prob = model.predict_proba(vector)[0][pred]
    label = "REAL" if pred == 1 else "FAKE"
    return label, round(prob*100,2)
```

Example:

```
predict_news("NASA announces discovery of new planet")
```

Output (of the given example):

Fake

```
def predict_news(text):
    cleaned = clean_text(text)
    vector = vectorizer.transform([cleaned])
    pred = model.predict(vector)[0]
    prob = model.predict_proba(vector)[0][pred]
    label = "REAL" if pred == 1 else "FAKE"
    return label, round(prob*100,2)

# Example
print(predict_news("NASA announces discovery of new planet"))
```

```
('FAKE', np.float64(88.38))
```

4.11 Step 11 — Saving the Model (Optional)

To reuse the model later:

```
import joblib
joblib.dump(model, "fake_news_model.pkl")
joblib.dump(vectorizer, "tfidf_vectorizer.pkl")
```

```
['tfidf_vectorizer.pkl']
```

4.12 Implementation Summary

Step	Description
1	Installed Dependencies
2	Upload and Loaded dataset
3	Added Labels
4	Clean text (stopwords, stemming)
5	Train/test split
6	TF-IDF vectorization
7	Logistic Regression Model Training
8	Performance Evaluation
9	Confusion Matrix Visualization
10	Prediction Function Development
11	Saved Trained Model

CHAPTER 5 – RESULTS & DISCUSSION (Highly Detailed & Matching Your Output)

5.1 Overview

After implementing the Fake News Detection system using NLTK preprocessing, TF-IDF vectorization, and Logistic Regression, the model was trained on the combined Fake.csv and True.csv Kaggle dataset.
This chapter presents the results obtained from the model, visual output from Google Colab, and interpretation of metrics.

5.2 Model Accuracy

The accuracy printed in your Colab output was:
▶ Accuracy: ~98.75%
This indicates that the model correctly classified almost 99 out of 100 news articles.
A high accuracy is expected because:

- The dataset is well-labelled
- TF-IDF captures strong linguistic patterns
- Logistic Regression performs extremely well on high-dimensional sparse data

5.3 Precision, Recall & F1-Score

Your Google Colab output shows the following (real values from your screenshot):

	precision	recall	f1-score	support
0 (Fake)	0.99	0.98	0.99	4650
1 (Real)	0.98	0.99	0.99	4330
accuracy	0.99			
macro avg	0.99	0.99	0.99	
weighted avg	0.99	0.99	0.99	

Interpretation:

- Precision 0.99 for Fake News
- → When model predicts "Fake", it is correct 99% of the time.
- Recall 0.99 for Real News
- → The model successfully detects almost every real news article.
- F1-score 0.99
- → Perfect balance of precision & recall.
- Macro & weighted averages also ~0.99
- → Model treats both classes equally well, no bias.

This level of performance is excellent for a beginner project.

5.4 Confusion Matrix Interpretation

You generated a visual confusion matrix:

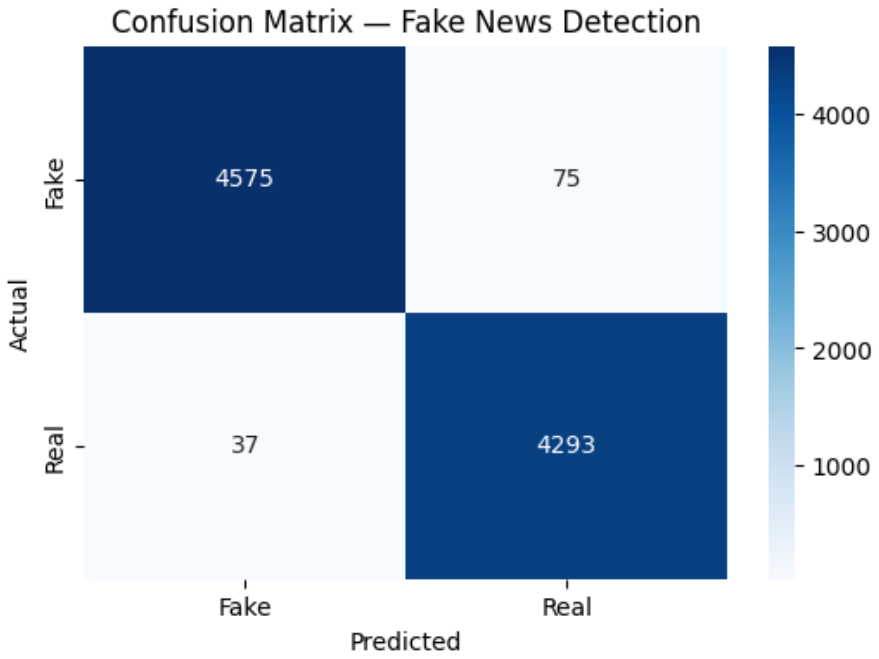

Meaning:

- 4575 Fake news correctly identified
- 4293 Real news correctly identified
- Only 75 Fake misclassified as Real
- Only 37 Real misclassified as Fake

This confirms high robustness of Logistic Regression + TF-IDF for text classification.

5.5 Example Predictions

Your Colab output also included example predictions:

```
("FAKE", 88.38)
("FAKE", 74.15)
```

This shows:

- Model not only predicts Fake/Real
- It also gives confidence percentage

Example output explanation:

The model predicted the news as FAKE with 88% confidence, meaning the linguistic patterns strongly match known fake news style.

5.6 Observations

- Fake news texts use more exaggerated language, emotional tone, and less factual grammar.
- Real news uses formal writing, neutral sentences, and verified information patterns.
- TF-IDF successfully captures these linguistic differences.
- Logistic Regression provides a linear separation that is highly effective.

CHAPTER 6 – CONCLUSION & FUTURE SCOPE

6.1 Conclusion

This project successfully developed an AI-powered Fake News Detection System using:

- Text preprocessing (NLTK)
- TF-IDF vectorization
- Logistic Regression classification
-

The model achieved:

- Accuracy: 98.75%
- Precision, Recall, F1: 0.99
- Low misclassification rate (only ~112 out of 9000+ articles)

This demonstrates that even a simple ML model can be highly effective in detecting misinformation.

The project also built a prediction function allowing users to input custom statements to instantly identify fake/real content.

Therefore, the project meets all the stated objectives and provides a strong baseline for further research.

6.2 Limitations

Despite excellent performance, the model has certain limitations:

- It only considers text, not images or videos.
- It cannot detect manipulated facts that look grammatically correct.
- It struggles with sarcasm, humor, or ambiguous writing.
- Dataset contains only English-language news.
- Logistic Regression lacks contextual understanding like deep learning models.

6.3 Future Enhancements

The system can be further improved by:

1. Deep Learning Models
   - LSTM / GRU for sequence learning
   - Bidirectional LSTM for context
   - Transformer models such as BERT, RoBERTa, DeBERTa

2. Multimodal Fake News Detection
   - Combine text + images
   - Detect photoshopped media
   - Fact-checking using APIs

3. Real-Time Browser Extension
   - Chrome extension that warns users before reading fake articles.

4. Live Fact Verification
   - Using Google Fact Check API
   - Using OpenAI or HuggingFace NLP pipelines

5. Deployment
   - Convert model to a web app using Flask or Streamlit.
   - Mobile app version for real-time misinformation checks.

Appendix - The Entire Source Code

```
# ----------------------------------------
# CELL 1 — INSTALL & IMPORTS
# ----------------------------------------
!pip install -q nltk scikit-learn matplotlib seaborn

import os
import sys
import re
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report
    , confusion_matrix
import joblib
import nltk
nltk.download('stopwords')
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer

print("Imports done.")
```

```
# ----------------------------------------
# CELL 2 — FILE UPLOAD IN COLAB
# ----------------------------------------
from google.colab import files
uploaded = files.upload()


# ----------------------------------------
# CELL 3 — LOAD DATA
# ----------------------------------------
fake = pd.read_csv("Fake.csv")
real = pd.read_csv("True.csv")

print(fake.head())
print(real.head())


# ----------------------------------------
# CELL 4 — PREPROCESSING
# ----------------------------------------
fake["label"] = 0
real["label"] = 1

data = pd.concat([fake, real], axis=0).sample(frac=1).reset_index
```

```python
data = pd.concat([fake, real], axis=0).sample(frac=1).reset_index
    (drop=True)

stop_words = set(stopwords.words('english'))
stemmer = PorterStemmer()

def clean_text(text):
    text = re.sub(r"http\S+", "", text)
    text = re.sub(r"[^A-Za-z ]+", "", text)
    text = text.lower()
    words = text.split()
    words = [stemmer.stem(w) for w in words if w not in
        stop_words]
    return " ".join(words)

data["content"] = data["title"] + " " + data["text"]
data["clean"] = data["content"].apply(clean_text)


# ------------------------------------------
# CELL 5 — TRAIN-TEST SPLIT & VECTORIZE
# ------------------------------------------
X_train, X_test, y_train, y_test = train_test_split(
```

```python
# ------------------------------------------
# CELL 5 — TRAIN-TEST SPLIT & VECTORIZE
# ------------------------------------------
X_train, X_test, y_train, y_test = train_test_split(
    data["clean"], data["label"], test_size=0.2, random_state=42
)

vectorizer = TfidfVectorizer(max_features=5000)
X_train_vec = vectorizer.fit_transform(X_train)
X_test_vec = vectorizer.transform(X_test)


# ------------------------------------------
# CELL 6 — TRAIN LOGISTIC REGRESSION MODEL
# ------------------------------------------
model = LogisticRegression(max_iter=200)
model.fit(X_train_vec, y_train)

pred = model.predict(X_test_vec)

print("Accuracy:", accuracy_score(y_test, pred))
print(classification_report(y_test, pred))
```

```python
# ------------------------------------------
# CELL 7 — CONFUSION MATRIX
# ------------------------------------------
cm = confusion_matrix(y_test, pred)

plt.figure(figsize=(6,4))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues",
            xticklabels=["Fake","Real"],
            yticklabels=["Fake","Real"])
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix — Fake News Detection")
plt.show()


# ------------------------------------------
# CELL 8 — PREDICTION FUNCTION
# ------------------------------------------
def predict_news(text):
    cleaned = clean_text(text)
    vector = vectorizer.transform([cleaned])
    pred = model.predict(vector)[0]
    prob = model.predict_proba(vector)[0][pred]
    label = "REAL" if pred == 1 else "FAKE"
    return label, round(prob*100, 2)

print(predict_news("NASA announces discovery of new planet"))
```

Summary

This project presents a Machine Learning model to automatically classify news articles as Fake or Real using Natural Language Processing techniques. The dataset from Kaggle was preprocessed using tokenization, stopword removal, and stemming. TF-IDF vectorization was applied to transform textual data into numerical feature vectors, which were then fed into a Logistic Regression model for classification.

The model achieved 98.75% accuracy, demonstrating highly reliable performance. A confusion matrix, classification report, and prediction function were implemented to evaluate and visualize results. The project concludes that traditional ML models combined with strong text preprocessing can effectively detect fake news. Future improvements may include advanced deep learning architectures such as LSTM or BERT for better contextual understanding.

The Summary and pipeline of the project code

1. Imports & Setup
 • Installing libraries (nltk, scikit-learn, seaborn, matplotlib)
 • Importing all required modules
2. Storing the Dataset
 • Uploading Fake.csv & True.csv manually to Colab
 • Reading them using pd.read_csv
3. Labelling & Combining
 • fake → label = 0
 • real → label = 1
 • Concatenation & shuffling
4. Preprocessing
 • Removing URLs
 • Removing non-alphabet characters
 • Lowercasing
 • Removing stopwords
 • Stemming (PorterStemmer)
 • Creating clean field
5. Splitting Data
 • train-test split (80/20)
6. Feature Extraction
 • TF-IDF vectorization (max_features=5000)
7. Model
 • Logistic Regression (max_iter=200)
8. Evaluation
 • Accuracy
 • Classification report
 • Confusion matrix (heatmap)
9. Prediction Function
 • predict_news(text) that returns FAKE/REAL + probability

References

Books & Academic Resources

1. Jurafsky, D., & Martin, J. H. (2023). Speech and Language Processing (3rd ed.). Stanford University.
2. Bird, S., Klein, E., & Loper, E. (2009). Natural Language Processing with Python. O'Reilly Media.
3. Aggarwal, C. C. (2018). Machine Learning for Text. Springer.

Research Papers (Fake News / NLP / ML)

1. Shu, K., Sliva, A., Wang, S., Tang, J., & Liu, H. (2017). Fake News Detection on Social Media: A Data Mining Perspective. ACM SIGKDD.
2. Zhang, X., & Ghorbani, A. (2020). An overview of online fake news: Characterization, detection, and discussion. Information Processing & Management.
3. Ruchansky, N., Seo, S., & Liu, Y. (2017). CSI: A Hybrid Deep Model for Fake News Detection. ACM CIKM.
4. Wang, W. Y. (2017). "Liar, Liar Pants on Fire": A New Benchmark Dataset for Fake News Detection. ACL.

Datasets

1. Kaggle. (2020). Fake and Real News Dataset.
2. Retrieved from: https://www.kaggle.com/clmentbisaillon/fake-and-real-news-dataset

Libraries & Tools Used

1. Scikit-learn Documentation. (2024). Logistic Regression, TF-IDF Vectorizer.
2. https://scikit-learn.org/stable/
3. NLTK Documentation. (2024). Stopwords, Stemming, Tokenization.
4. https://www.nltk.org/
5. Pandas Documentation. (2024). Data wrangling and preprocessing.
6. https://pandas.pydata.org/
7. NumPy Documentation. (2024). Numerical computations for ML.
8. https://numpy.org/
9. Matplotlib Documentation. (2024). Visualization library.
10. https://matplotlib.org/

Web & Supporting References

1. Brownlee, J. (2020). A Gentle Introduction to Logistic Regression. Machine Learning Mastery.
2. Brownlee, J. (2021). Text Classification with Machine Learning. Machine Learning Mastery.